



EvoMind-*V* · Single-Image Architecture

Frozen SigLIP2 vision tower + trainable projector + selected 768-dim CISPO text backbone.


(a) Image features



(b) Token Alignment



(c) Language Model



Gray: frozen parameters Purple: updated parameters / learned features

Current training: 2 epochs · batch 4 · BF16 · LR 5e-6. No SigLIP text tower, cross-attention module or video encoder is used.